

Netflix Movie Recommendation Web application

Cloud Computing Final Project Fall 2020

Ruinan Zhang rz1109
Jiahui Wang jw6273
Wei Si ws1623

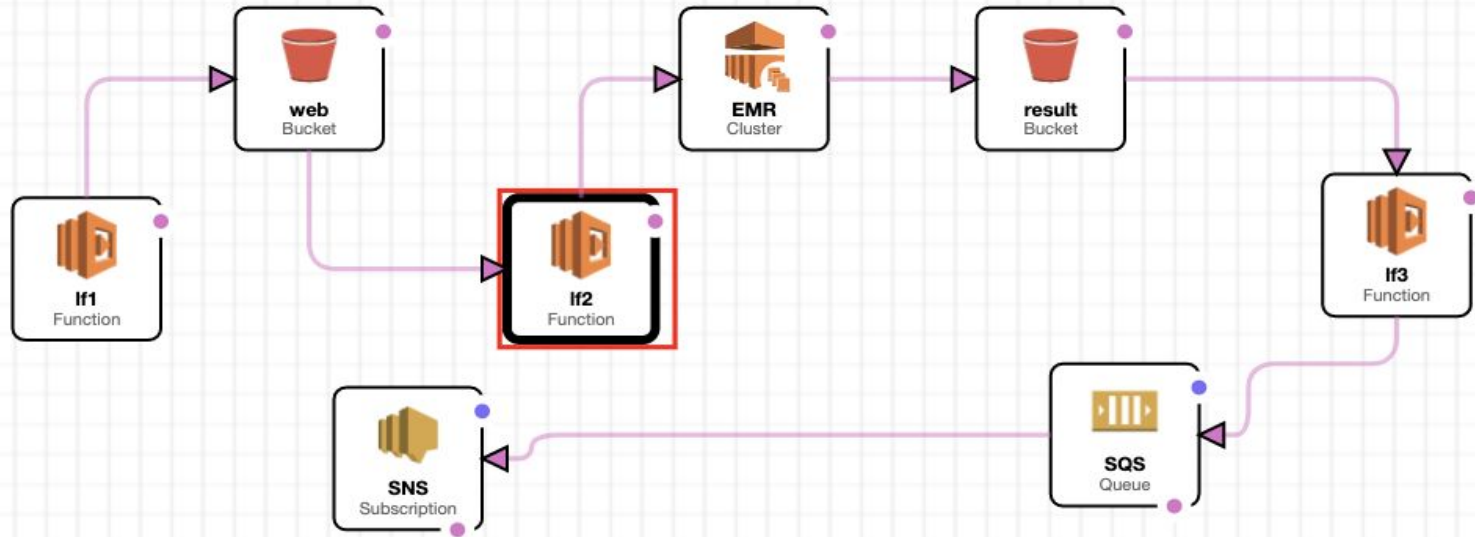
Overview

The project will provide Netflix movie recommendations to the user based on their personalized ratings for movies. Users will be able to select movies they like from the front-end website and then the system will provide customized recommendations for the users using Alternating Least Squares algorithms with the help of PySpark ML. The result of recommendation will be sent to users via text message after the computation is done.

Project Architecture

- Amazon S3 for storage of data and program for cluster and host front-end website
- Amazon EC2 instance for submitting Apache Spark jobs
- Amazon EMR cluster(3 nodes cluster) with Spark 2.4.7 for running Spark jobs
- Alternating Least Squares algorithm (implemented in Apache Spark ML that built for a large-scale collaborative filtering problems) for recommendation
- DynamoDB for movie details storage and user info storage
- API gateway for connecting front-end to the back-end
- Lambda functions to get query and for each query submit different spark jobs and retrieve data and send them via SNS
- Amazon SNS service for delivering results to the users

Project Structure



Dataset- Netflix Prize data

(<https://www.kaggle.com/netflix-inc/netflix-prize-data>)

The original movie rating files contain over 100 million ratings from 480 thousand randomly-chosen, anonymous Netflix customers over 17 thousand movie titles. The data were collected between October, 1998 and December, 2005 and reflect the distribution of all ratings received during this period. The ratings are on a scale from 1 to 5 (integral) stars.

Number of Users: 750

Number of Movies: 1,000

Number of Ratings: 420,000

Input Ratings Data File contains (Before Cleaning): movie_id, user_id, ratings, date_of_rating

Input Movie Title file contains: movie_id, year_of_release, movie_title

Note: Due to the time and budget limit, we will be using 50,000 ratings instead of 4,20,000 in this demo project.**

Sample data (after cleaning)

Movie rating data:

Movie_id, user_id, rating, rating date

```
49995      26    2459417  4    2004-08-14
49996      26    1469394  4    2004-07-18
49997      26    302426   3    2004-08-27
49998      26    1455293  1    2004-07-20
49999      26    2076552  4    2004-09-14
50000      26    976191   5    2004-08-24
```

Movie title csv:

Movie_id, year, movie_title

```
movie_titles.csv
1    1,2003,Dinosaur Planet
2    2,2004,Isle of Man TT 2004 Review
3    3,1997,Character
4    4,1994,Paula Abdul's Get Up & Dance
5    5,2004,The Rise and Fall of ECW
6    6,1997,Sick
7    7,1992,8 Man
8    8,2004,What the #$*! Do We Know!?
9    9,1991,Class of Nuke 'Em High 2
10   10,2001,Fighter
```

Front-end website:

<http://cc-final-proj-web.s3-website-us-west-2.amazonaws.com/>

- Use questionnaire to collect user's rating on 5 movies. (We planned to use more data but then computation would take too much time, so we decided to do the computation based on 5 ratings, using 26 movies in total)
- Collect user name and phone number for sending out recommendation results
- The data collected from the front-end would be sent to Lambda Function 'proj_insert_rating_py2' via API 'movie-user-input'

Get your personalized movie recommendation

Take this movie rating questionnaire to find out!

Your name:

Your phone:

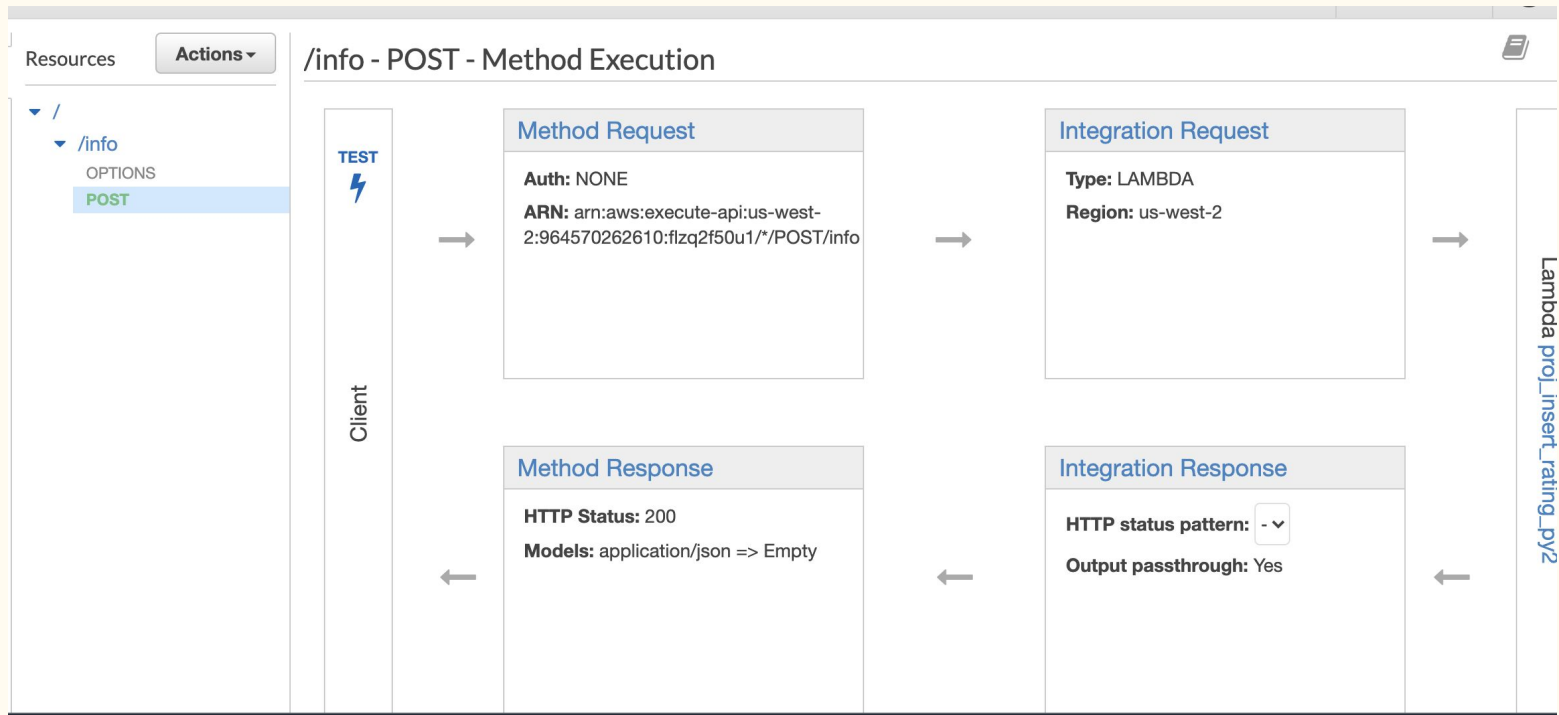
How do you like 'Dinosaur Planet'?

- ☐ Really love it! (☆☆☆☆)
- ☐ Like it! (☆☆☆☆)
- ☒ I think it's okay. (☆☆☆)
- ☐ Not really my type. (☆☆)
- ☐ I don't like it :((☆)

How do you like 'My Bloody Valentine'?

- ☐ Really love it! (☆☆☆☆)
- ☐ Like it! (☆☆☆☆)

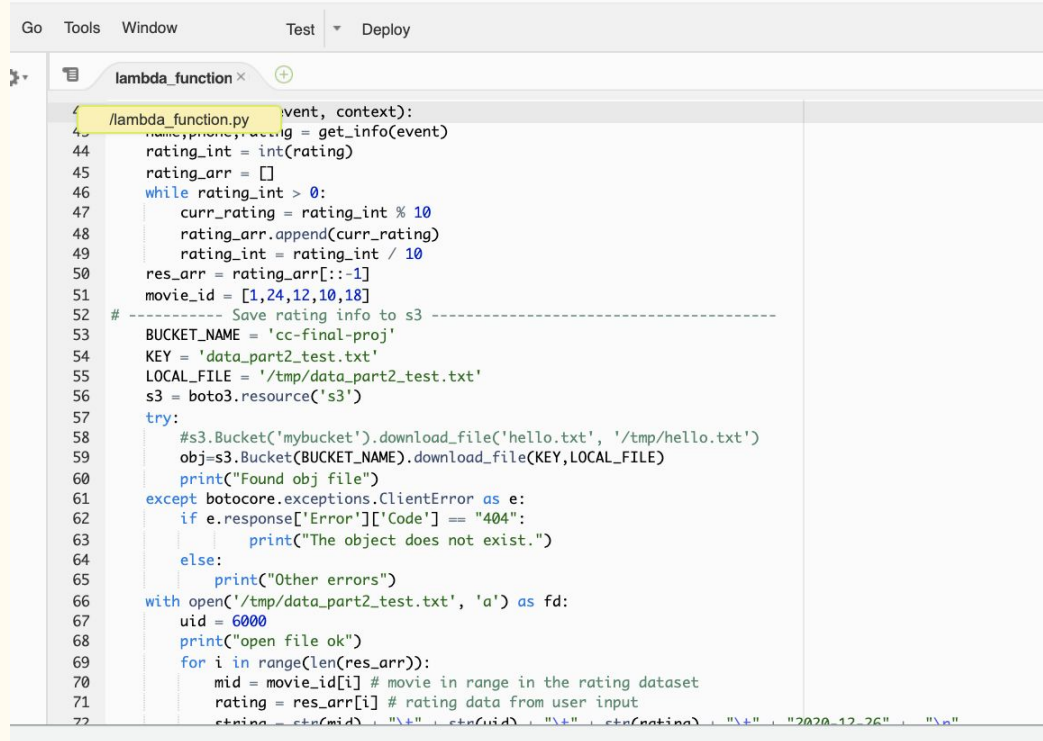
API 'movie-user-input'



Lambda Function 1 - 'proj_insert_rating_py2'

The LF1 will be triggered by API Gateway:

movie-user-input and then save the newly collected user rating info to existing using rating data set. At the same time, it will save user's info and phone number to DynamoDB for the further usage after getting the recommendation results.



```
Go Tools Window Test Deploy

lambda_function x

43 /lambda_function.py def handler(event, context):
44     name, phone, rating = get_info(event)
45     rating_int = int(rating)
46     rating_arr = []
47     while rating_int > 0:
48         curr_rating = rating_int % 10
49         rating_arr.append(curr_rating)
50         rating_int = rating_int // 10
51     res_arr = rating_arr[::-1]
52     movie_id = [1,24,12,10,18]
53     # ----- Save rating info to s3 -----
54     BUCKET_NAME = 'cc-final-proj'
55     KEY = 'data_part2_test.txt'
56     LOCAL_FILE = '/tmp/data_part2_test.txt'
57     s3 = boto3.resource('s3')
58     try:
59         #s3.Bucket('mybucket').download_file('hello.txt', '/tmp/hello.txt')
60         obj=s3.Bucket(BUCKET_NAME).download_file(KEY,LOCAL_FILE)
61         print("Found obj file")
62     except botocore.exceptions.ClientError as e:
63         if e.response['Error']['Code'] == "404":
64             print("The object does not exist.")
65         else:
66             print("Other errors")
67     with open('/tmp/data_part2_test.txt', 'a') as fd:
68         uid = 6000
69         print("open file ok")
70         for i in range(len(res_arr)):
71             mid = movie_id[i] # movie in range in the rating dataset
72             rating = res_arr[i] # rating data from user input
73             string = str(mid) + "\t" + str(uid) + "\t" + str(rating) + "\t" + "2020-12-26" + "\n"
```

Lambda Function 2- 'submit-pyspark-job'

The LF2 will be triggered by S3 bucket: cc-final-proj.

Procedure:

- 1, User submits a rating data from front end to S3 bucket
- 2, Trigger the LF2
- 3, LF2 submits a pyspark job to the AWS EMR cluster

AWS EMR (Elastic Map Reducer)

The goal we use Amazon EMR is to simplify the big data structure. The EMR cluster we build here includes three ec2 instances:

1 Master node and 2 Core nodes.

Instances (3)

Info

↻

Connect

Instance state ▾

Actions ▾

Launch instances

▾

🔍 Filter instances

< 1 >

⚙

Instance state: running ✕

Clear filters

| <input type="checkbox"/> | Name ▾ | Instance ID | Instance state ▾ | Instance type ▾ | Status check | A.. | Availability Zone |
|--------------------------|-----------------|-------------------------------------|---|-----------------|---------------------------------------|-----|-------------------|
| <input type="checkbox"/> | cc-final-proj | i-01bcbe082849042c5 | <div><div>✔</div>Running<div>🔍🔊</div></div> | m5.xlarge | <div><div>✔</div>2/2 checks ...</div> | + | us-west-2d |
| <input type="checkbox"/> | cc-final-proj-2 | i-05de549eb77b4b969 | <div><div>✔</div>Running<div>🔍🔊</div></div> | m5.xlarge | <div><div>✔</div>2/2 checks ...</div> | + | us-west-2d |
| <input type="checkbox"/> | cc-final-proj-3 | i-01b61e5a3f139c222 | <div><div>✔</div>Running<div>🔍🔊</div></div> | m5.xlarge | <div><div>✔</div>2/2 checks ...</div> | + | us-west-2d |

Recommendation Function: ALSForUid.py

This function runs in the EMR cluster. It takes the user id and rating data as input. Then use the ALS algorithm to get the most related ten movies and save the result into the S3 bucket.

Procedure:

- 1, Input the user's uid and rating
- 2, Train the data using ALS model
- 3, Output the result into S3 bucket: cc-final-result


Reference: <https://spark.apache.org/docs/2.2.0/ml-collaborative-filtering.html>

```
20
21
22 f getRecommendationsFromALSLibrary(uid):
23
24     '''
25     Build the recommendation model using Alternating Least Squares
26     '''
27
28     parts = ratings_input.map(lambda row: row.value.split("\t"))
29     ratingsRDD = parts.map(lambda p: Row(userId=int(p[1]), movieId=int(p[0]), rating=float(p[2])))
30     ratings = spark.createDataFrame(ratingsRDD)
31
32
33     #Need a training and test set into 85 and 15%
34     train, test = ratings.randomSplit([0.85,0.15],123)
35
36     print ("Training Count for Ratings", train.count())
37     print ("Test Count for Ratings",test.count())
38
39     # Cache Training and Test Data
40     train.cache()
41     test.cache()
42
43     # Generate ALS Model.
44     als = ALS(maxIter=5, regParam=0.01, userCol="userId", itemCol="movieId", ratingCol="rating", coldStartStrategy="drop")
45     model = als.fit(train)
46
47     # Evaluate the model by computing the RMSE on the test data
48     predictions = model.transform(test)
49     evaluator = RegressionEvaluator(metricName="rmse", labelCol="rating",predictionCol="prediction")
50     rmse = evaluator.evaluate(predictions)
51     print("Root Mean Square Error = " + str(rmse))
52
53     # Generate Top 10 movie recommendations for each user
54     userRecs = model.recommendForAllUsers(10)    ## <class 'pyspark.sql.dataframe.DataFrame'>
55
56     print ("Printing Recommendations for All Users: ")
57     res = userRecs.filter(userRecs.userId == uid) #dataframe
58     rec_df = res.withColumn("recommendations", explode("recommendations")).select("*, col("recommendations")["movieId"].alias("movieId"), col("recommendations")["rating"].alias("
59     rec_df = rec_df.drop('recommendations') # drop this col
60
61     print ("Printing Recommendations for User ", uid)
62     print(rec_df.show())
63     return rec_df
```

ALS Algorithm

ALS: Alternating Least Squares

| | Movie 1 | Movie 2 | Movie ... | Movie N |
|----------|---------|---------|-----------|---------|
| User 1 | 1 | BLANK | BLANK | 3 |
| User 2 | BLANK | 5 | BLANK | 3 |
| User 3 | BLANK | BLANK | 1 | BLANK |
| User 4 | 2 | 3 | BLANK | BLANK |
| User 5 | BLANK | BLANK | 1 | BLANK |
| User 6 | 4 | BLANK | 5 | BLANK |
| User 7 | BLANK | 4 | BLANK | BLANK |
| User ... | BLANK | 3 | BLANK | BLANK |
| User m | BLANK | BLANK | BLANK | 4 |



| | Movie 1 | Movie 2 | Movie ... | Movie N |
|----------|---------|---------|-----------|---------|
| User 1 | 1 | 4 | 2 | 3 |
| User 2 | 1 | 5 | 3 | 3 |
| User 3 | 2.5 | 2.8 | 1 | 3.5 |
| User 4 | 2 | 3 | 2 | 3.5 |
| User 5 | 2.5 | 2.8 | 1 | 3.5 |
| User 6 | 4 | 1.2 | 5 | 1.4 |
| User 7 | 1 | 4 | 2.5 | 3 |
| User ... | 2 | 3 | 2 | 3 |
| User m | 1 | 4 | 2 | 4 |

$$P_u := P_u - \alpha \cdot \frac{\partial L}{\partial P_u} = P_u - \alpha \cdot [\sum_i 2(P_u^T Q_i - R_{ui})Q_i + 2\lambda P_u]$$

$$Q_i := Q_i - \alpha \cdot \frac{\partial L}{\partial Q_i} = Q_i - \alpha \cdot [\sum_u 2(P_u^T Q_i - R_{ui})P_u + 2\lambda Q_i]$$

Spark Job Results

Objects (2)
Objects are the fundamental entities stored in Amazon S3. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)



↻

Delete

Actions ▾

Create folder

Upload

| | Name | Type | Last modified |
|-------------------------------------|--|------|---|
| <input checked="" type="checkbox"/> |  part-00000-fcb80d92-c8ee-4be9-8355-1ce10be3409b-c000.csv | csv | December 27, 2020, 17:41:29 (UTC-05:00) |
| <input type="checkbox"/> |  _SUCCESS | - | December 27, 2020, 17:41:29 (UTC-05:00) |

Spark job's results will return recommendation for the user, and then we use the movieID to locate more detailed info about the movie in DynamoDB

```
Users > kathola > Downloads > part-0000
1  |userId,movieId,rating
2  |6000,10,4.993286
3  |6000,12,3.950684
4  |6000,23,3.2662668
5  |6000,20,3.1579885
6  |6000,4,2.358785
7  |6000,14,2.3062017
8  |6000,6,1.9251564
9  |6000,5,1.6568274
10 |6000,7,1.2981827
11 |6000,1,1.1842637
12
```

Create item

Actions ▾



Scan: [Table] movie_titles: movieId ^

Scan ▾

[Table] movie_titles: movieId ▾ ^

+ Add filter

Start search

| <input type="checkbox"/> | movieId ⓘ | title | yearOfRelease |
|--------------------------|-----------|---|---------------|
| <input type="checkbox"/> | 1 | Dinosaur Planet | 2003 |
| <input type="checkbox"/> | 10 | Fighter | 2001 |
| <input type="checkbox"/> | 11 | Full Frame: Documentary Shorts | 1999 |
| <input type="checkbox"/> | 12 | My Favorite Brunette | 1947 |
| <input type="checkbox"/> | 13 | Lord of the Rings: The Return of the King: Extended Edition: Bonus Material | 2003 |
| <input type="checkbox"/> | 14 | Nature: Antarctica | 1982 |
| <input type="checkbox"/> | 15 | Neil Diamond: Greatest Hits Live | 1988 |
| <input type="checkbox"/> | 16 | Screamers | 1996 |
| <input type="checkbox"/> | 17 | 7 Seconds   | 2005 |

Lambda Function 3- 'proj_send_rec_py2'

This function is triggered by the S3 bucket cc-final-result.

Procedure:

- 1, S3 bucket cc-final-result receives the updating recommendation result
- 2, Trigger the LF3
- 3, LF3 submit the sending message request to AWS SQS.
- 4, SQS triggers the AWS SNS service and sends text message to user.

Individual Contributions

Ruinan Zhang: Front-end website / APIGateway / LF1

Wei Si: Spark job Submission / LF2 / EMR cluster creation

Jiahui Wang: ALS algorithm development / LF3 / SNS service